

MIT 5.2 - Senior IT Project

Phase 1: Product Definition Document

DocVerify AI

AI-Powered Document Verification System

Winter Semester 2026/27

Lecturer: Dr. Ali Alhosseini

Submitted by	Sachin Adihalli Ravigowda
Student ID	111121997
University email	saravigowda@stud.mediadesign.de
Submission date	01.11.2026
GitHub Project (backlog)	https://github.com/users/SachinAR-21/projects/1
GitHub repository	https://github.com/SachinAR-21/DocVerify-AI

Table of Contents

1. Motivation . . . 3

1.1. The problem with manual verification . . . 3

1.2. Why AI is a suitable solution . . . 3

2. Product Vision . . . 5

3. Personas . . . 6

3.1. Priya Raman - Applicant / Submitter . . . 6

3.2. Jana Becker - Reviewer / Verifier . . . 7

3.3. Tobias Klein - Administrator . . . 8

4. Scenarios . . . 9

4.1. Priya uploads her documents for a Master's application . . . 9

4.2. Jana reviews AI verification reports on a Monday morning . . . 9

4.3. Tobias updates the rules for a new intake . . . 9

5. User Stories and Acceptance Criteria . . . 11

5.1. Applicant / Submitter . . . 11

5.2. Reviewer / Verifier . . . 13

5.3. Administrator . . . 16

6. Prioritized Feature List (MoSCoW) . . . 19

6.1. Reasons for the “Won't” decisions . . . 20

7. Mapping of Features to User Stories . . . 21

8. Product Backlog (GitHub Project) . . . 23

1. Motivation

Almost every organisation that admits, hires or onboards people has to check documents. A university looks at transcripts and language certificates, an employer checks degrees and work permits, a bank or a landlord checks ID cards, bank statements and proof of address. In most places this is still done by hand: someone opens the PDF, compares it with a rule sheet or an old example, checks a name against what was typed in a form and makes a call. It works, but it does not scale, and it is a lot less reliable than people like to admit.

1.1. The problem with manual verification

It costs a lot of time. A careful check of one application with four or five documents easily takes 15 to 25 minutes. During peak season an admissions office may receive several thousand applications within a few weeks. The result is a backlog, overtime for staff and applicants who wait weeks for an answer that often depends on a single unreadable scan.

It is error-prone and inconsistent. Reviewers are human. After the fortieth transcript of the day, small things slip through: a date of birth that does not match the passport, an expired language certificate, a name spelled two different ways. Different reviewers also interpret the rules differently, so the same document can be accepted on Monday and rejected on Thursday.

It is risky. Forged and manipulated documents are getting easier to make. Free editing tools can change a grade or a date in minutes and the result often looks clean to the naked eye. Accepting a fake document can mean an unqualified student, a fraudulent loan or a legal problem for the organisation. On top of that, documents are often sent around by email and stored in shared folders, which is a real data protection problem under the GDPR.

It is frustrating for everyone. Applicants do not know where their file stands and often find out weeks later that one page was unreadable. Reviewers spend most of their time on routine checks instead of the cases that actually need their judgement.

1.2. Why AI is a suitable solution

Most of the work in document verification is reading, comparing and checking against rules. This is exactly where current AI techniques are strong:

- **OCR and PDF text extraction** turn scans and photos into machine-readable text.
- **Large Language Models (LLMs)** can classify the document type and pull out structured fields such as name, date of birth, institution, grades or amounts, even when every institution uses a different layout.

- **Retrieval-Augmented Generation (RAG)** lets the system check each document against the organisation's own rules, reference templates, grade conversion tables and lists of recognised institutions, and cite exactly which rule it used.
- **Automated consistency checks** compare every field across all documents of an application, every time, without getting tired.
- **Authenticity indicators** such as mismatched fonts, impossible dates or inconsistent totals can be flagged for a human to look at.

We are deliberately **not** proposing to replace the reviewer. AI makes mistakes too, and decisions about people's education, jobs or housing should not be fully automated. DocVerify AI does the repetitive groundwork, explains what it found and where that rule comes from, and a human makes the final decision. Our goal is to cut the time per application roughly in half, make decisions more consistent and give applicants faster and clearer feedback.

2. Product Vision

> *"For **universities, employers, public offices, banks and landlords** who **need to verify large numbers of applicant documents quickly, reliably and in a GDPR-compliant way**, DocVerify AI is an **AI-powered document verification platform** that **automatically extracts, cross-checks and verifies documents against official rules and produces an explainable report, so that reviewers can make confident decisions in minutes instead of hours**. Unlike **manual checking via email and shared folders or generic OCR tools**, our product **explains every finding with citations to the rules it used, keeps a human in control of every decision, and gives applicants instant feedback on their uploads**."*

In plain words: checking a document should feel less like detective work and more like confirming a well-prepared summary. Applicants should know within seconds whether their upload is usable, reviewers should see at a glance what needs their attention and why, and administrators should be able to prove at any time who decided what, based on which rule, model and prompt version.

Guiding principles

- **Human in the loop:** the AI recommends, people decide.
- **Explainable, not a black box:** every finding comes with a reason and a cited source.
- **Traceable:** every AI decision is logged with the prompt version, model and sources used.
- **Privacy by design:** encrypted storage, role-based access, automatic deletion, and only synthetic or anonymised data during development.
- **Simple for applicants:** uploading should work from a phone without instructions.

3. Personas

The personas below represent the three roles defined for the platform: applicant, reviewer and administrator. They are fictional, but each one is put together from the kind of people and complaints we expect to meet in a typical university admissions setting.

3.1. Priya Raman - Applicant / Submitter

Age	26
Location	Chennai, India
Situation	Applying for a Master's in Data Science at three German universities
Devices	Android phone (main device), shared family laptop
Quote	“I just want to know if they got my documents and if something is wrong.”

Priya finished her Bachelor's in computer engineering two years ago and has been working as a junior developer since. Applying abroad means collecting her degree certificate, transcripts, passport, IELTS result and proof of financial resources, then uploading them to several different portals. Last year one of her applications was rejected because a transcript page was “unreadable”, and she only found out after the deadline.

Goals

- Submit a complete, correct application before the deadline.
- Know immediately if a document has a problem so she can fix it.
- Track the status of her application without writing emails.
- Be sure her passport and personal data are handled safely.

Frustrations

- Portals that accept any file and only complain weeks later.
- No status updates; emails to the admissions office go unanswered.
- Having to re-upload the whole application because of one bad file.
- Not understanding why a document was rejected.

Technical skills

Comfortable with smartphones and web apps, uses scanner apps on her phone. Not familiar with PDF tools or file formats beyond the basics. Her internet connection at home is sometimes slow.

3.2. Jana Becker - Reviewer / Verifier

Age	41
Location	Berlin, Germany
Role	Admissions officer, International Office of a mid-sized university
Workload	Around 50-70 applications per day in peak season
Quote	“I'm not worried about the normal cases. I'm worried about the one fake I don't catch.”

Jana has worked in admissions for twelve years. She knows the usual documents from dozens of countries and can often tell when something looks off. Her day is mostly spent opening attachments, comparing names and dates with the application form, looking up grade conversion tables and writing the same emails asking for better scans. She cares about fairness and does not want a machine making decisions about applicants on its own.

Goals

- Get through her queue without overtime during application season.
- Spend her time on unclear or suspicious cases, not routine checks.
- Make consistent decisions that she can justify with the actual rule.
- Have the document, the extracted data and the AI report on one screen.

Frustrations

- Switching between email, shared drives, rule PDFs and the campus management system.
- Bad scans that cost time before she can even start checking.
- Forgeries that are getting harder to spot by eye.
- Tools that give a result but no reason (“black boxes”).

Technical skills

Confident office user (Outlook, Excel, the university's campus management system). Not a technical person and has no interest in configuring AI models; she wants a clear interface that just works and answers she can trust.

3.3. Tobias Klein - Administrator

Age	34
Location	Berlin, Germany
Role	IT system administrator for the university's student services
Responsibility	Users, document types, verification rules, knowledge base, data protection
Quote	“If the data protection officer asks me who saw a passport, I need an answer in five minutes.”

Tobias looks after several systems used by student services. He sets up accounts for new staff, keeps the admission rules up to date and works closely with the data protection officer. He is the person who gets the call when something breaks, when a rule changes or when an audit is coming up.

Goals

- Control exactly who can see and change what.
- Update document types, rules and knowledge-base documents without waiting for a developer.
- Know which prompt version and model produced each AI decision.
- Have a complete audit trail and make sure personal data is deleted on time.

Frustrations

- Tools without proper roles where everyone is effectively an admin.
- Rules hidden in old PDFs that nobody knows are still valid.
- Personal data scattered across mailboxes and network drives.
- Systems that cannot export logs or statistics.

Technical skills

High. Comfortable with user and role management, REST APIs, logs and basic scripting. Understands GDPR requirements in practice and has a basic understanding of how LLMs work.

4. Scenarios

4.1. Priya uploads her documents for a Master's application

It is 21:30 and the application deadline is tomorrow. Priya registers on DocVerify AI with her email address, creates a new application for the Master's in Data Science and opens it on her phone. She selects “Degree certificate” as the document type and takes a picture. After a few seconds the app tells her the photo is too dark and the bottom edge is cut off. She moves to the window, takes the photo again and this time gets a green check mark. Her transcripts are already PDFs on her laptop, so she adds them there and links them to the same application. When she uploads her IELTS certificate she adds a note that her surname is misspelled as “Ramen” on the certificate. Before submitting, she sees a checklist confirming that all required documents are present.

Four days later she gets an email: her transcript needs a certified English translation. She logs in, sees exactly which document is affected and why, uploads the translation and does not have to touch the rest of her application. A week later the status changes to “Verified”.

4.2. Jana reviews AI verification reports on a Monday morning

On Monday morning Jana has 63 new applications in her queue. DocVerify AI processed them over the weekend: it extracted the data, checked it against the admission rules in the knowledge base and created a verification report for each one. 48 are marked **Verified** with a high confidence score. She opens a few, looks at the extracted fields next to the original documents and the list of passed checks, and accepts them in under two minutes each.

One application is marked **Rejected**. The report highlights the grade table of a transcript, explains that the font in two cells differs from the rest of the document and that the final average does not match the listed grades, and cites the template for that university. Jana zooms in, agrees, and requests a certified copy sent directly by the issuing university. Another file is marked **Needs Review** because of an IELTS score of 6.0. She asks the AI assistant: “Does this applicant meet the language requirements?” The assistant answers that the programme requires 6.5 and cites the admission rules document, but also notes that the applicant uploaded a newer TOEFL result. Jana checks it, overrides the decision with a short justification and moves on. By lunch she has cleared the queue.

4.3. Tobias updates the rules for a new intake

The university has changed its language requirement for one programme from IELTS 6.5 to 7.0 and is opening a new programme that also requires a portfolio. Tobias logs into the admin area, uploads the new admission rules PDF to the knowledge base and marks the old version as archived, so the AI only cites the current rule. He creates the new document type

“Portfolio”, marks it as required for the new programme, and gives two student assistants the “Reviewer (read-only)” role for peak season.

A week later the data protection officer asks which staff accessed a specific applicant's passport, and why one application was rejected. Tobias searches the audit log by application ID, exports every view, download and decision with timestamps, and shows that the rejection was produced with prompt version v1.3, which rules were cited, and which reviewer confirmed it. He also checks that the retention rule is active: documents of rejected applicants are deleted automatically six months after the decision.

5. User Stories and Acceptance Criteria

All user stories follow the format “As a [role], I want [goal], so that [benefit]”. Each story has acceptance criteria that tell us when it is actually done, and a MoSCoW priority. All stories are also managed as issues in the GitHub Project (see section 8).

Role	Number of stories	Story IDs
Applicant	8	US-01, US-02, US-03, US-04, US-05, US-18, US-19, US-20
Reviewer	9	US-06, US-07, US-08, US-09, US-10, US-11, US-21, US-22, US-25
Administrator	9	US-12, US-13, US-14, US-15, US-16, US-17, US-23, US-24, US-26

5.1. Applicant / Submitter

US-01 – Priority: Must

As an **applicant**, I want to upload my documents as PDF, JPG or PNG from my phone or computer, so that I can submit my application without needing a scanner.

Acceptance criteria:

- Files can be added via drag and drop, file picker or phone camera.
- PDF, JPG and PNG up to 10 MB per file are accepted; other formats show a clear error.
- Upload progress is shown and an interrupted upload can be retried.
- Files are transferred over HTTPS and stored encrypted.

US-02 – Priority: Should

As an **applicant**, I want to get instant feedback if my upload is blurry, too dark or cut off, so that I can fix it right away instead of being rejected weeks later.

Acceptance criteria:

- A readability check runs within 10 seconds after upload.
- Problems are explained in plain language (e.g. “The bottom edge is cut off”).
- The applicant can replace the file directly from the message.
- The applicant can still submit with a warning if they are sure the file is fine.

US-03 – Priority: Must

As an **applicant**, I want to see the status of each of my documents, so that I know where my application stands without having to write emails.

Acceptance criteria:

- Each document shows one of: Uploaded, Processing, Verified, Needs Review, Rejected.
- A checklist shows which required documents are still missing.
- The status page works on mobile.

US-04 – Priority: Must

As an **applicant**, I want to be notified with a clear reason when a document is not accepted, so that I can replace just that document before the deadline.

Acceptance criteria:

- An email notification is sent when a document status changes to Needs Review or Rejected, or when the reviewer requests more documents.
- The notification includes the reason written by the reviewer.
- Only the affected document needs to be re-uploaded; the rest of the application stays unchanged.

US-05 – Priority: Must

As an **applicant**, I want to understand how my data is used and to request its deletion, so that I can trust the system with my passport and personal information.

Acceptance criteria:

- Before the first upload, a short privacy notice is shown and consent is recorded.
- The applicant can download a copy of their stored data.
- A deletion request is logged and handled according to the retention policy.

US-18 – Priority: Could

As an **applicant**, I want to use the portal and submit documents in English or German, so that I understand every instruction and my documents are processed correctly.

Acceptance criteria:

- The user interface is available in English and German.
- German and English documents are extracted and verified correctly.
- Notifications are sent in the selected language.

US-19 – Priority: Must

As an **applicant**, I want to register and log in with my email address and a password, so that my documents and application are only accessible to me.

Acceptance criteria:

- Registration requires a valid email address and a password meeting the password policy.
- The email address must be confirmed before the first upload.
- Passwords are stored hashed; after repeated failed logins the account is temporarily locked.

US-20 – Priority: Must

As an **applicant**, I want to select the document type, add notes and link several documents to one application, so that the reviewer sees my complete application together with my explanations.

Acceptance criteria:

- The applicant chooses a document type from a list before uploading.
- Optional notes can be added to each document.
- Several documents (e.g. degree, transcript, language certificate) are grouped under one application.
- The system warns if the detected document type differs from the selected one.

5.2. Reviewer / Verifier

US-06 – Priority: Must

As a **reviewer**, I want to see the data the AI extracted next to the original document, so that I can check the most important fields quickly without retyping anything.

Acceptance criteria:

- Extracted fields (name, date of birth, institution, issue/expiry date, document number, grades or amounts) are shown side by side with the document.
- Clicking a field highlights where it was found in the document.
- The reviewer can correct a wrongly extracted value; the correction is stored and logged.

US-07 – Priority: Must

As a **reviewer**, I want to see a confidence score and the reasons why a document was flagged, so that I can focus on suspicious documents and understand what to look at.

Acceptance criteria:

- Each document gets a confidence score and an overall result.
- Every flag (e.g. mismatched fonts, impossible dates, inconsistent totals) has a short explanation and highlights the affected area.
- The AI result is labelled as a recommendation, not a decision.

US-08 – Priority: Must

As a **reviewer**, I want a work queue that I can sort and filter by result, deadline and programme, so that I deal with the most urgent and risky applications first.

Acceptance criteria:

- The queue can be sorted by AI result, confidence, submission date and deadline.
- Filters exist for programme, status and assigned reviewer.
- An application opened by one reviewer is shown as “in progress” to others.

US-09 – Priority: Must

As a **reviewer**, I want to accept, reject or request additional documents with a comment, so that I can make and document the final decision in one place.

Acceptance criteria:

- Accept, Reject and Request documents are available on the application view.
- Reject and Request documents require a comment that is sent to the applicant.
- Every decision is stored with reviewer name and timestamp.

US-10 – Priority: Must

As a **reviewer**, I want to override an AI decision and give a justification, so that wrong results do not block applicants and the prompts can be improved later.

Acceptance criteria:

- The reviewer can override the AI result and must enter a justification.
- Overrides and corrections are logged in the audit trail.
- Overrides can be exported to evaluate and improve prompts.

US-11 – Priority: Must

As a **reviewer**, I want the system to check consistency across all documents of an application, so that I immediately see mismatches in names, dates of birth or other data.

Acceptance criteria:

- Name, date of birth and nationality are compared across all documents and the application form.
- Mismatches are highlighted; small differences (date formats, accents, transliteration) are tolerated.
- The reviewer can see the conflicting values side by side.

US-21 – Priority: Must

As a **reviewer**, I want an explainable verification report for each document and application, so that I understand the AI result and can trust or challenge it.

Acceptance criteria:

- The report shows the overall result: Verified, Needs Review or Rejected.
- It lists passed and failed checks (completeness, validity, consistency, compliance, authenticity).
- Each finding has a natural-language explanation and a citation to the rule or knowledge-base source used.

US-22 – Priority: Should

As a **reviewer**, I want to ask an AI assistant questions about a document or application, so that I get quick answers without searching through rule documents.

Acceptance criteria:

- The reviewer can type questions such as “Does this applicant meet the language requirements?”
- Answers are based only on the document data and the knowledge base and cite their sources.
- If the information is missing, the assistant says so instead of guessing.

US-25 – Priority: Could

As a **reviewer**, I want to export a verification report as PDF, so that I can attach it to the applicant's file or share it with colleagues.

Acceptance criteria:

- A “Download PDF” button is available on every report.
- The PDF contains result, checks, explanations, citations and reviewer decision.
- The export is logged in the audit trail.

5.3. Administrator

US-12 – Priority: Must

As an **administrator**, I want to create users and assign roles with defined permissions, so that everyone only sees and does what they are allowed to.

Acceptance criteria:

- Roles include Applicant, Reviewer, Reviewer (read-only) and Administrator.
- Only administrators can change roles.
- Access to a page without permission is denied and logged.

US-13 – Priority: Must

As an **administrator**, I want to configure document types and the verification rules applied to them, so that new programmes and requirements can be supported without code changes.

Acceptance criteria:

- Admins can add, edit and deactivate document types.
- For each type, required/optional status and the fields to extract can be set.
- Rules such as minimum scores or maximum document age can be defined per programme.

US-14 – Priority: Must

As an **administrator**, I want a complete audit log of all actions and AI decisions, so that I can answer data protection and compliance questions quickly.

Acceptance criteria:

- Uploads, views, downloads, edits, decisions and deletions are logged with user and time.
- Every AI decision is logged with prompt version, model and sources used.
- The log can be filtered and exported as CSV and cannot be edited by any user.

US-15 – Priority: Should

As an **administrator**, I want a dashboard with verification numbers and AI quality, so that I can see bottlenecks and whether the AI is actually helping.

Acceptance criteria:

- Shows number of verifications, results, and average processing time per week.
- Shows AI accuracy and the reviewer override rate.
- Data can be filtered by programme and time period.

US-16 – Priority: Must

As an **administrator**, I want to set retention periods after which documents are deleted automatically, so that we follow the GDPR without manual clean-up.

Acceptance criteria:

- Retention period can be set per decision outcome (e.g. rejected: 6 months).
- Deletion runs automatically and is recorded in the audit log.
- Admins get a summary of deletions each month.

US-17 – Priority: Could

As an **administrator**, I want a public REST API to connect DocVerify AI to other systems, so that verified results flow into our existing processes without copy-and-paste.

Acceptance criteria:

- Endpoints exist for creating an application, uploading documents and reading results.
- The API is secured with tokens and documented with OpenAPI.
- Calls are rate-limited and logged.

US-23 – Priority: Must

As an **administrator**, I want to upload and update knowledge-base documents such as rules, templates and grade tables, so that the AI always verifies against the current official rules.

Acceptance criteria:

- Admins can upload, replace and archive knowledge-base documents (PDF, DOCX, TXT).
- New documents are automatically split, embedded and made searchable.
- Archived documents are no longer used or cited in new verifications.

US-24 – Priority: Should

As an **administrator**, I want to manage and version the prompts used by the AI, so that changes to the AI behaviour are controlled and can be rolled back.

Acceptance criteria:

- Prompts are stored as versioned templates, not hard-coded.
- Admins can view, compare and activate a prompt version.
- Each AI result records which prompt version was used.

US-26 – Priority: Could

As an **administrator**, I want to upload and verify documents in bulk, so that large batches, for example from a partner university, can be processed at once.

Acceptance criteria:

- Several files or a ZIP archive can be uploaded in one step.
- Progress of the batch is shown and failed files are listed.
- Results appear in the reviewer queue like normal applications.

6. Prioritized Feature List (MoSCoW)

We derived the features from the user stories and prioritised them with the MoSCoW method. **Must** features are needed for a usable, safe and explainable first release, including the core AI pipeline. **Should** features add a lot of value and are planned directly after. **Could** features are nice to have if time allows. **Won't** features are consciously left out of this project.

In total there are 27 features: 15 Must, 4 Should, 4 Could and 4 Won't.

ID	Feature	Priority	User stories
F-01	User registration, login and authentication	Must	US-19
F-02	Role-based access control and user management	Must	US-12
F-03	Document upload with type selection, notes and file validation	Must	US-01, US-20
F-04	Application management (linking several documents)	Must	US-20, US-03
F-05	OCR / PDF text extraction and document classification	Must	US-06, US-20
F-06	LLM-based structured field extraction (JSON)	Must	US-06
F-07	RAG knowledge base (ingestion, embeddings, retrieval)	Must	US-23, US-21
F-08	AI verification checks (completeness, validity, consistency, compliance, authenticity)	Must	US-07, US-11, US-21
F-09	Explainable verification report with confidence and citations	Must	US-21, US-07
F-10	Reviewer dashboard and work queue	Must	US-08
F-11	Human review: accept, override, correct, request documents	Must	US-09, US-10, US-06
F-12	Applicant status tracking and notifications	Must	US-03, US-04
F-13	Document types and verification rule management	Must	US-13
F-14	Audit log and AI traceability	Must	US-14
F-15	GDPR consent, data export and retention rules	Must	US-05, US-16
F-16	Upload readability check (blur, lighting, cropping)	Should	US-02
F-17	AI assistant chat with cited answers	Should	US-22
F-18	Prompt management and versioning	Should	US-24
F-19	Analytics dashboard (volume, accuracy, time, override rate)	Should	US-15
F-20	Export of verification reports as PDF	Could	US-25
F-21	Batch upload and bulk verification	Could	US-26
F-22	Multi-language interface and documents (EN/DE)	Could	US-18

ID	Feature	Priority	User stories
F-23	Public REST API for third-party integration	Could	US-17
F-24	Fully automatic approval without human review	Won't	–
F-25	QR-code / digital signature verification	Won't	–
F-26	Live verification with issuing institutions' databases	Won't	–
F-27	Native iOS / Android app	Won't	–

6.1. Reasons for the “Won't” decisions

- **F-24 Fully automatic approval:** decisions about admission, jobs or housing affect people's lives. The AI supports reviewers but never decides alone, which also respects Art. 22 GDPR.
- **F-25 QR-code / digital signature verification:** listed as optional in the brief. Formats differ by issuer, so we postpone it until the core verification pipeline is stable.
- **F-26 Live verification with issuers:** there is no common interface for universities and test providers worldwide, so it is not realistic within this project.
- **F-27 Native apps:** a responsive web app with camera access covers the applicant's needs.

7. Mapping of Features to User Stories

The table maps every user story to the feature(s) that deliver it. Every user story is covered by at least one feature, and every feature we plan to build is backed by at least one user story.

Story	Role	Short description	Priority	Feature(s)
US-01	Applicant	To upload my documents as PDF, JPG or PNG from my phone or computer	Must	F-03
US-02	Applicant	To get instant feedback if my upload is blurry, too dark or cut off	Should	F-16
US-03	Applicant	To see the status of each of my documents	Must	F-04, F-12
US-04	Applicant	To be notified with a clear reason when a document is not accepted	Must	F-12
US-05	Applicant	To understand how my data is used and to request its deletion	Must	F-15
US-06	Reviewer	To see the data the AI extracted next to the original document	Must	F-05, F-06, F-11
US-07	Reviewer	To see a confidence score and the reasons why a document was flagged	Must	F-08, F-09
US-08	Reviewer	A work queue that I can sort and filter by result, deadline and programme	Must	F-10
US-09	Reviewer	To accept, reject or request additional documents with a comment	Must	F-11
US-10	Reviewer	To override an AI decision and give a justification	Must	F-11
US-11	Reviewer	The system to check consistency across all documents of an application	Must	F-08
US-12	Administrator	To create users and assign roles with defined permissions	Must	F-02
US-13	Administrator	To configure document types and the verification rules applied to them	Must	F-13
US-14	Administrator	A complete audit log of all actions and AI decisions	Must	F-14
US-15	Administrator	A dashboard with verification numbers and AI quality	Should	F-19
US-16	Administrator	To set retention periods after which documents are deleted automatically	Must	F-15

Story	Role	Short description	Priority	Feature(s)
US-17	Administrator	A public REST API to connect DocVerify AI to other systems	Could	F-23
US-18	Applicant	To use the portal and submit documents in English or German	Could	F-22
US-19	Applicant	To register and log in with my email address and a password	Must	F-01
US-20	Applicant	To select the document type, add notes and link several documents to one application	Must	F-03, F-04, F-05
US-21	Reviewer	An explainable verification report for each document and application	Must	F-07, F-08, F-09
US-22	Reviewer	To ask an AI assistant questions about a document or application	Should	F-17
US-23	Administrator	To upload and update knowledge-base documents such as rules, templates and grade tables	Must	F-07
US-24	Administrator	To manage and version the prompts used by the AI	Should	F-18
US-25	Reviewer	To export a verification report as PDF	Could	F-20
US-26	Administrator	To upload and verify documents in bulk	Could	F-21

8. Product Backlog (GitHub Project)

All 26 user stories are managed as issues in the GitHub repository and added to the GitHub Project board. Each issue contains the user story, its acceptance criteria as a checklist, a MoSCoW priority label and the related feature(s). New issues from the repository are added to the board automatically.

GitHub Project (backlog)	https://github.com/users/SachinAR-21/projects/1
GitHub repository	https://github.com/SachinAR-21/DocVerify-AI
Board columns	Todo, In Progress, Done
Labels	user story; priority: must / should / could